

CSE 417T: Homework 2

Hangxiao Zhu

October 7, 2022

Problem 1.

a. Report on each of the maximum number of iterations:

Maximum Iterations	Training time	Iterations	E_{in}	e on training set	e on test set
10000	0.07678604125976562	10000	0.5847165561996043	0.3092105263157895	0.31724137931034485
100000	0.7698559761047363	100000	0.4937022212468817	0.2236842105263158	0.20689655172413793
1000000	7.694957971572876	1000000	0.43535262292978577	0.1513157894736842	0.1310344827586207

Based on the binary classification errors in the training and test, the model performs no worse in the test set than in the training set. That means the generalization properties of the model is very strong.

By increasing the maximum number of iterations, the binary classification errors on test data gradually decreases, that is, the model's performance gets improved by increasing the number of training iterations.

b. I normalized both datasets using the training mean and variance.

The mean I used is [54.5855263, 0.657894737, 3.125, 130.532895, 249.835526, 0.131578947, 0.914473684, 149.565789, 0.269736842, 1.00263158, 1.57894737, 0.651315789, 4.80263158].

The variance I used is [78.8084747, 0.225069252, 0.938322368, 315.735760, 3199.82163, 0.114265928, 0.986106302, 550.258830, 0.196978878, 1.17551939, 0.362188366, 0.819208795, 3.76367729].

The standard deviation I used is [8.87741374, 0.47441464, 0.96867041, 17.76895495, 56.56696591, 0.33803244, 0.99302885, 23.45759642, 0.44382303, 1.08421372, 0.60182088, 0.90510154, 1.94001992].

Report on each of the learning rate:

η	Training time	Iterations	E_{in}	e on training set	e on test set
0.01	7.577642917633057	1000000	0.4073814412083029	0.17105263157894737	0.1103448275862069
0.1	7.762326002120972	1000000	0.4073814412083028	0.17105263157894737	0.1103448275862069
1	0.006412029266357422	819	0.40738144120830283	0.17105263157894737	0.1103448275862069
4	0.0015261173248291016	186	0.40738144120830283	0.17105263157894737	0.1103448275862069
7	0.001306772232055664	164	0.40738144120830283	0.17105263157894737	0.1103448275862069
7.5	0.005311012268066406	678	0.40738144120830294	0.17105263157894737	0.1103448275862069
7.6	0.023669004440307617	3051	0.40738144120830283	0.17105263157894737	0.1103448275862069
7.7	8.204280138015747	1000000	0.40843667430393826	0.16447368421052633	0.11724137931034483

Normalizing the data did affect the performace of the model. By observing the table above, we can see the cross-entropy error on the training data set, the binary classification error on both the training and test data sets all decreases after normalizing the data. Thus, I think normalizing the data boost the performace of the model.

Changing learning rate η affects the number of iterations the gradient descent takes to converge. Choosing a proper learning rate can significantly improve the efficiency of the model and reduce the number of iterations required to converge. However, if the learning rate is too small or too large, the number of iterations required to converge will be large.

Problem 2.

Based on $E_{out}(g^{(\mathcal{D})}) = \mathbb{E}_{\mathbf{x},y}[(g^{(\mathcal{D})}(\mathbf{x}) - y(x))^2]$ and $y(x) = f(x) + \epsilon$, we can deduce that:

$$\begin{aligned}
\mathbb{E}_{\mathcal{D}}[E_{out}(g^{(\mathcal{D})})] &= \mathbb{E}_{\mathcal{D}}[\mathbb{E}_{\mathbf{x},y}[(g^{(\mathcal{D})}(\mathbf{x}) - y(x))^2]] \\
&= \mathbb{E}_{\mathbf{x},y}[\mathbb{E}_{\mathcal{D}}[(g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}) + \bar{g}(\mathbf{x}) - y(x))^2]] \\
&= \mathbb{E}_{\mathbf{x},y}[\mathbb{E}_{\mathcal{D}}[(g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}))^2 + (\bar{g}(\mathbf{x}) - y(x))^2 + 2(g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}))(\bar{g}(\mathbf{x}) - y(x))]] \\
&\quad (\text{Note that } \mathbb{E}_{\mathcal{D}}[(g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}))(\bar{g}(\mathbf{x}) - y(x))] = (\bar{g}(\mathbf{x}) - y(x))\mathbb{E}_{\mathcal{D}}[(g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}))] = 0) \\
&= \mathbb{E}_{\mathbf{x},y}[\mathbb{E}_{\mathcal{D}}[(g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}))^2 + (\bar{g}(\mathbf{x}) - y(x))^2]] \\
&= \mathbb{E}_{\mathbf{x},y}[\mathbb{E}_{\mathcal{D}}[(g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}))^2]] + \mathbb{E}_{\mathbf{x},y}[(\bar{g}(\mathbf{x}) - y(x))^2] \\
&= \mathbb{E}_{\mathbf{x},y}[\text{Variance of } g^{(\mathcal{D})}(\mathbf{x})] + \mathbb{E}_{\mathbf{x},y}[(\bar{g}(\mathbf{x}) - f(x) - \epsilon)^2] \\
&= \mathbb{E}_{\mathbf{x},y}[\text{Variance of } g^{(\mathcal{D})}(\mathbf{x})] + \mathbb{E}_{\mathbf{x},y}[(\bar{g}(\mathbf{x}) - f(x))^2 - 2(\bar{g}(\mathbf{x}) - f(x))\epsilon + \epsilon^2] \\
&= \mathbb{E}_{\mathbf{x},y}[\text{Variance of } g^{(\mathcal{D})}(\mathbf{x})] + \mathbb{E}_{\mathbf{x},y}[(\bar{g}(\mathbf{x}) - f(x))^2] + \mathbb{E}_{\mathbf{x},y}[\epsilon^2] - 2\mathbb{E}_{\mathbf{x},y}[(\bar{g}(\mathbf{x}) - f(x))\epsilon] \\
&= \mathbb{E}_{\mathbf{x}}[\text{Variance of } g^{(\mathcal{D})}(\mathbf{x})] + \mathbb{E}_{\mathbf{x}}[\text{Bias of } \bar{g}(\mathbf{x})] + \mathbb{E}_{\mathbf{x},y}[\epsilon^2] - 2\mathbb{E}_{\mathbf{x},y}[(\bar{g}(\mathbf{x}) - f(x))\epsilon] \\
&= \text{var} + \text{bias} + \mathbb{E}_{\mathbf{x},y}[\epsilon^2] - 2\mathbb{E}_{\mathbf{x},y}[(\bar{g}(\mathbf{x}) - f(x))\epsilon] \\
&= \text{var} + \text{bias} + \mathbb{E}_{\epsilon}[\epsilon^2] - 2\mathbb{E}_{\mathbf{x}}[(\bar{g}(\mathbf{x}) - f(x))\epsilon] \\
&= \text{var} + \text{bias} + \mathbb{E}_{\epsilon}[\epsilon^2] - 2\mathbb{E}_{\mathbf{x}}[(\bar{g}(\mathbf{x}) - f(x))]\mathbb{E}_{\mathbf{x}}[\epsilon] \\
&= \text{var} + \text{bias} + \mathbb{E}_{\epsilon}[\epsilon^2] - 2\mathbb{E}_{\mathbf{x}}[(\bar{g}(\mathbf{x}) - f(x))]\mathbb{E}_{\epsilon}[\epsilon] \\
&\quad (\text{Note that } \epsilon \text{ is a zero mean noise, which means } \bar{\epsilon} = \mathbb{E}_{\epsilon}[\epsilon] = 0) \\
&= \text{var} + \text{bias} + \mathbb{E}_{\epsilon}[(\epsilon - \bar{\epsilon})^2] \\
&= \text{var} + \text{bias} + \sigma^2 \\
&= \sigma^2 + \text{bias} + \text{var}
\end{aligned}$$

Problem 3.

(a) Since the input variable x is uniformly distributed in the interval $[-1, 1]$, we know $\bar{x} = \mathbb{E}_{\mathbf{D}}[x] = 0$. Based on them, we can deduce that:

$$\begin{aligned}
\bar{g}(x) &= \mathbb{E}_{\mathbf{D}}[g(x)] \\
&= \mathbb{E}_{\mathbf{D}}[ax + b] \\
&\quad (\text{Note that } a = \frac{x_2^2 - x_1^2}{x_2 - x_1} = x_1 + x_2, \quad b = \frac{x_1x_2^2 - x_2x_1^2}{x_2 - x_1} = -x_1x_2) \\
&= \mathbb{E}_{\mathbf{D}}[(x_1 + x_2)x - x_1x_2] \\
&= (\mathbb{E}_{\mathbf{D}}[x_1] + \mathbb{E}_{\mathbf{D}}[x_2])x - \mathbb{E}_{\mathbf{D}}[x_1]\mathbb{E}_{\mathbf{D}}[x_2] \\
&= 0
\end{aligned}$$

(b) Experiment to determine $\bar{g}(x)$:

- Sample a x from the interval $[-1, 1]$ according to uniform distribution.
- Iterate t times (i.g. $t = 1000$):
 - Sample two datapoints uniformly from the interval $[-1, 1]$.
 - Compute a and b .
 - Compute the value of $g^{(\mathcal{D})}(x)$ and save the result in an array.
- Compute $\bar{g}(x)$ by taking the average of the array of $g^{(\mathcal{D})}(x)$.

Experiment to determine E_{out} , bias , var :

- Iterate t times (i.g. $t = 10000$):
 - Sample a x from the interval $[-1, 1]$ according to uniform distribution.
 - Follow the experiment described above to calculate the array of $g^{(\mathcal{D})}(x)$ and $\bar{g}(x)$.
 - Use $\mathbb{E}_{\mathcal{D}}[(g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}))^2]$ to calculate var and save the result in an array.
 - Use $(\bar{g}(\mathbf{x}) - f(\mathbf{x}))^2$ to calculate $bias$ and save the result in an array.
 - Use $\mathbb{E}_{\mathcal{D}}[(g^{(\mathcal{D})}(\mathbf{x}) - f(\mathbf{x}))^2]$ to calculate E_{out} and save the result in an array.
- Compute E_{out} , $bias$, var by taking the average of the arrays of E_{out} , $bias$, var .

(c) As shown in the table below, $E_{out} \approx bias + var$, which is consistent with the theoretical result.

E_{out}	$bias$	var	$bias + var$
0.5430268888689767	0.2051878468504339	0.33732433111079485	0.5425121779612287

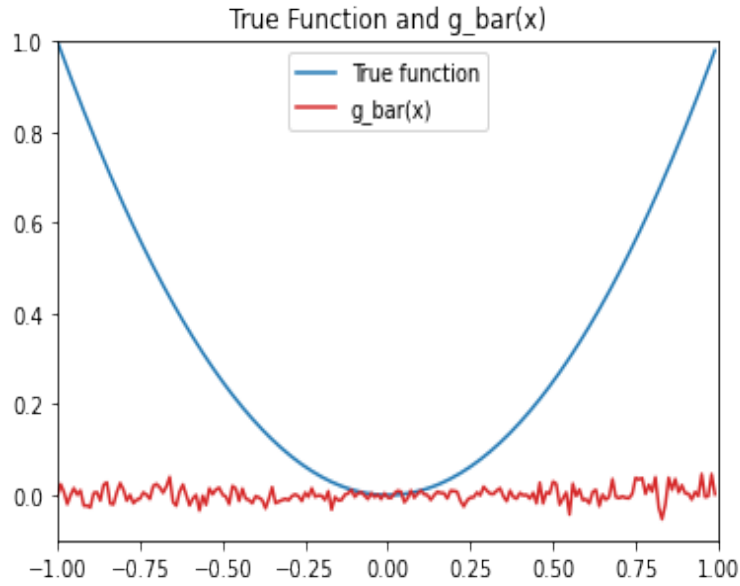


Figure 1: Histogram of Number of Iterations

(d) Calculate $bias$:

$$\begin{aligned}
 bias &= \mathbb{E}_{\mathbf{x}}[(\bar{g}(\mathbf{x}) - f(\mathbf{x}))^2] \\
 &= \mathbb{E}_{\mathbf{x}}[f(\mathbf{x})^2] \\
 &= \mathbb{E}_{\mathbf{x}}[(x^2)^2] \\
 &= \mathbb{E}_{\mathbf{x}}[x^4] \\
 &= \frac{1}{2} \int_{-1}^1 x^4 dx \\
 &= \frac{1}{5}
 \end{aligned}$$

Calculate var :

$$\begin{aligned}
var &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}_{\mathcal{D}}[(g^{(\mathcal{D})}(\mathbf{x}) - \bar{g}(\mathbf{x}))^2]] \\
&= \mathbb{E}_{\mathbf{x}}[\mathbb{E}_{\mathcal{D}}[(a\mathbf{x} + b)^2]] \\
&\quad (\text{Note that } a = \frac{x_2^2 - x_1^2}{x_2 - x_1} = x_1 + x_2, \quad b = \frac{x_1x_2^2 - x_2x_1^2}{x_2 - x_1} = -x_1x_2) \\
&= \mathbb{E}_{\mathbf{x}}[\mathbb{E}_{\mathcal{D}}[(x_1 + x_2)^2\mathbf{x}^2 + x_1^2x_2^2 - 2x_1x_2(x_1 + x_2)\mathbf{x}]] \\
&= \mathbb{E}_{\mathbf{x}}[\mathbb{E}_{\mathcal{D}}[x_1^2 + x_2^2 + 2x_1x_2]\mathbf{x}^2 + \mathbb{E}_{\mathcal{D}}[x_1^2x_2^2] - 2\mathbb{E}_{\mathcal{D}}[x_1^2x_2 + x_2^2x_1]\mathbf{x}] \\
&= \mathbb{E}_{\mathbf{x}}[(\mathbb{E}_{\mathcal{D}}[x_1^2] + \mathbb{E}_{\mathcal{D}}[x_2^2] + 2\mathbb{E}_{\mathcal{D}}[x_1x_2])\mathbf{x}^2 + \mathbb{E}_{\mathcal{D}}[x_1^2]\mathbb{E}_{\mathcal{D}}[x_2^2] - 2\mathbb{E}_{\mathcal{D}}[x_1x_2]\mathbb{E}_{\mathcal{D}}[x_1 + x_2]\mathbf{x}] \\
&\quad (\text{Note that } \mathbb{E}_{\mathcal{D}}[x^2] = \frac{1}{2} \int_{-1}^1 x^2 dx = \frac{1}{3}, \mathbb{E}_{\mathcal{D}}[x_1x_2] = 0) \\
&= \mathbb{E}_{\mathbf{x}}[\frac{2}{3}\mathbf{x}^2 + \frac{1}{9}] \\
&= \frac{2}{3}\mathbb{E}_{\mathbf{x}}[\mathbf{x}^2] + \frac{1}{9} \\
&= \frac{2}{3} \times \frac{1}{3} + \frac{1}{9} \\
&= \frac{1}{3}
\end{aligned}$$

Calculate E_{out} :

$$\begin{aligned}
\mathbb{E}_{\mathcal{D}}[E_{out}(g^{(\mathcal{D})})] &= \mathbb{E}_{\mathcal{D}}[\mathbb{E}_{out}[(g^{(\mathcal{D})}(\mathbf{x}) - y(x))^2]] \\
&= bias + var \\
&= \frac{1}{5} + \frac{1}{3} \\
&= \frac{8}{15}
\end{aligned}$$

Problem 4.

(a) We can write $E_n(\mathbf{w})$ as

$$E_n(\mathbf{w}) = \begin{cases} 0 & y_n \mathbf{w}^T \mathbf{x}_n \geq 1 \\ (1 - y_n \mathbf{w}^T \mathbf{x}_n)^2 & y_n \mathbf{w}^T \mathbf{x}_n < 1 \end{cases}$$

We can see that when $y_n \mathbf{w}^T \mathbf{x}_n > 1$, $E_n(\mathbf{w}) = 0$ is continuous and differentiable, when $y_n \mathbf{w}^T \mathbf{x}_n \rightarrow 1$, $E_n(\mathbf{w}) = 0$, when $y_n \mathbf{w}^T \mathbf{x}_n < 1$, $E_n(\mathbf{w}) = (1 - y_n \mathbf{w}^T \mathbf{x}_n)^2$ is continuous and differentiable.

The gradient $\nabla E_n(\mathbf{w})$ is

$$\nabla E_n(\mathbf{w}) = \begin{cases} 0 & y_n \mathbf{w}^T \mathbf{x}_n \geq 1 \\ -2y_n x_n (1 - y_n \mathbf{w}^T \mathbf{x}_n) & y_n \mathbf{w}^T \mathbf{x}_n < 1 \end{cases}$$

(b) I will prove by two cases:

Case 1: If $\text{sign}(\mathbf{w}^T \mathbf{x}_n) \neq y_n$, then $\mathbb{1}[\text{sign}(\mathbf{w}^T \mathbf{x}_n) \neq y_n] = 1$ and $y_n \mathbf{w}^T \mathbf{x}_n \leq 0$. Therefore, $E_n(\mathbf{w}) = (1 - y_n \mathbf{w}^T \mathbf{x}_n)^2 \geq 1$. Thus, $E_n(\mathbf{w}) \geq \mathbb{1}[\text{sign}(\mathbf{w}^T \mathbf{x}_n) \neq y_n]$.

Case 2: If $\text{sign}(\mathbf{w}^T \mathbf{x}_n) = y_n$, then $\mathbb{1}[\text{sign}(\mathbf{w}^T \mathbf{x}_n) \neq y_n] = 0$ and $y_n \mathbf{w}^T \mathbf{x}_n \geq 0$. Therefore, when $0 \leq y_n \mathbf{w}^T \mathbf{x}_n < 1$, $E_n(\mathbf{w}) = (1 - y_n \mathbf{w}^T \mathbf{x}_n)^2 \geq 0$, when $y_n \mathbf{w}^T \mathbf{x}_n \geq 1$, $E_n(\mathbf{w}) = 0$. In both situations, $E_n(\mathbf{w}) \geq \mathbb{1}[\text{sign}(\mathbf{w}^T \mathbf{x}_n) \neq y_n]$.

In conclusion, $E_n(\mathbf{w}) \geq \mathbb{1}[\text{sign}(\mathbf{w}^T \mathbf{x}_n) \neq y_n]$, which means $E_n(\mathbf{w})$ is an upper bound for $\mathbb{1}[\text{sign}(\mathbf{w}^T \mathbf{x}_n) \neq y_n]$. Hence, $\frac{1}{N} \sum_{n=1}^N E_n(\mathbf{w})$ is an upper bound for $E_{in}(\mathbf{w})$.

(c) The stochastic gradient descent should be in the following steps:

- Initialize a \mathbf{w} .

- Iterate t times:
 - Randomly choose x_n, y_n .
 - Let $s_t = \mathbf{w}^T \mathbf{x}_n$, then $y_n s_t = y_n \mathbf{w}^T \mathbf{x}_n$.
 - If $y_n \mathbf{w}^T \mathbf{x}_n \geq 1$, $\nabla E_n(\mathbf{w}) = 0$, then $w \leftarrow w - \eta \nabla E_n(\mathbf{w}) = w - 0 = w$.
 - If $y_n \mathbf{w}^T \mathbf{x}_n < 1$, $\nabla E_n(\mathbf{w}) = -2y_n x_n (1 - y_n \mathbf{w}^T \mathbf{x}_n)$, then $w \leftarrow w - \eta \nabla E_n(\mathbf{w}) = w + 2\eta y_n x_n (1 - y_n s_t)$. Then replace 2η by η' , $w = w + \eta' y_n x_n (1 - y_n s_t)$.

We can see the process above is the same as the Adaline algorithm, therefore the Adaline algorithm performs stochastic gradient descent on $\frac{1}{N} \sum_{n=1}^N E_n(\mathbf{w})$.

Problem 5.

- This transformation $\Phi(\mathbf{x})$ converts each data point into N dimension, although PLA has $d_{vc} = N + 1$, it is still very expensive to compute on this dataset. Besides, since the $d_{vc} = N + 1$, this hypothesis has a very large d_{vc} , which makes it hard to minimize the generalization error.
- This transformation $\phi_n(\mathbf{x})$ requires extra storage for each data which is not in the dataset. Therefore, more storage is required.
- When N gets large and the original data set has large dimensions, the storage will be significantly large because each data point will be converted into $101 - d$ data. Also, the computation will also be slow since the transformation requires to traverse i, j .